



Git Workflows & Continuous Integration

Better Scientific Software Tutorial

Jared O'Neal

Argonne National Laboratory

ISC High Performance Conference

June 16, 2019



See slide 2 for
license details

License, citation, and acknowledgments



License and Citation

- This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/) (CC BY 4.0).
- Requested citation: **Jared O'Neal, Git Workflows & Continuous Integration, in Better Scientific Software Tutorial, ISC High Performance Conference, Frankfurt, Germany, 2019. DOI: [10.6084/m9.figshare.8242859](https://doi.org/10.6084/m9.figshare.8242859)**

Acknowledgements

- This work was supported by the U.S. Department of Energy Office of Science, Office of Advanced Scientific Computing Research (ASCR), and by the Exascale Computing Project (17-SC-20-SC), a collaborative effort of the U.S. Department of Energy Office of Science and the National Nuclear Security Administration.
- This work was performed in part at the Argonne National Laboratory, which is managed by UChicago Argonne, LLC for the U.S. Department of Energy under Contract No. DE-AC02-06CH11357
- Anshu Dubey, Klaus Weide, Saurabh Chawdhary, and Carlo Graziani
- Iulian Grindeanu
- Alicia Klinvex



Git Workflows

Goals

Development teams would like to use version control to collaborate productively and ensure correct code

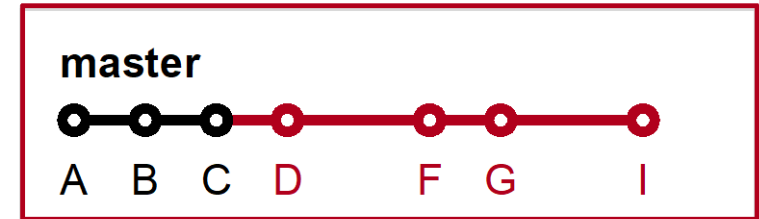
- Understand challenges related to parallel code development *via* distributed version control
- Understand extra dimensions of distributed version control & how to use them
 - Local vs. remote repositories
 - Branches
 - Issues, Pull Requests, & Code Reviews
- Exposure to workflows of different complexity
- What to think about when evaluating different workflows
- Motivate continuous integration

Distributed Version Control System (DVCS)

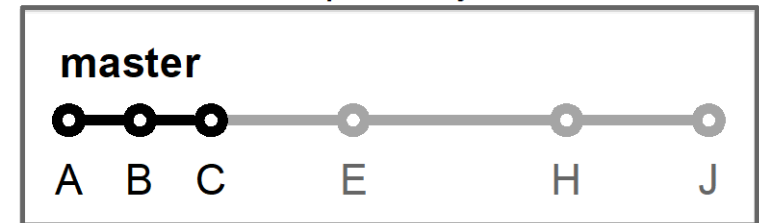
Two developers collaborating *via* Git

- Local copies of master branch synched to origin
- Each develops on **local** copy of master branch
- All copies of master immediately diverge
- How to **integrate** work on origin?

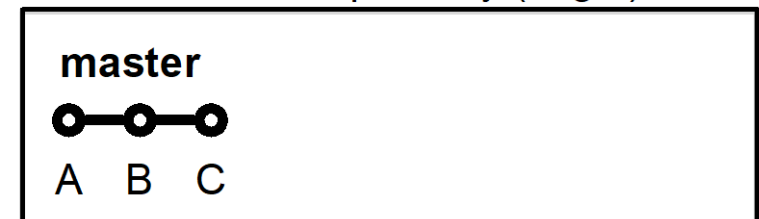
Alice's Local Repository



Bob's Local Repository



Main Remote Repository (origin)



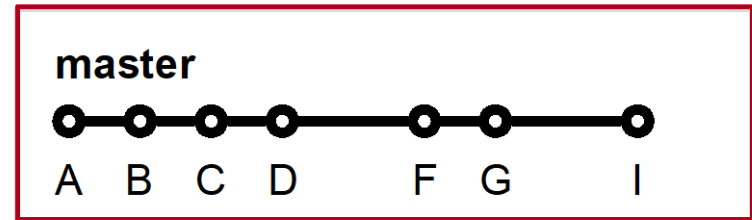
● = commit — = branch
X = commit ID

DVCS Race Condition

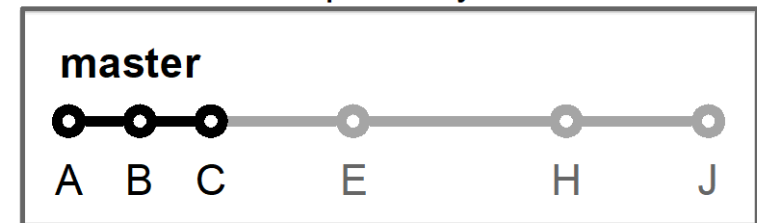
Integration of independent work occurs when local repos interact with remote repo

- Alice pushes her local commits to remote repo first
- No integration conflicts
- No risk
- Alice's local repo identical to remote repo

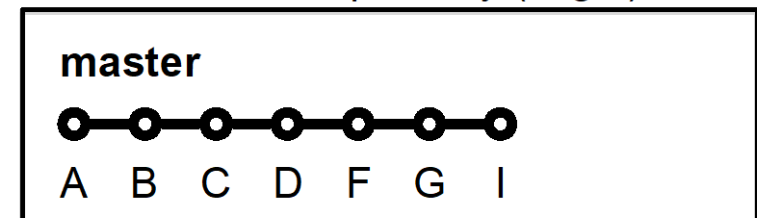
Alice's Local Repository



Bob's Local Repository



Main Remote Repository (origin)



● = commit — = branch
X = commit ID

Integration Conflicts Happen

Bob's push to remote repo is rejected

- Alice updated code in commit D
- Bob updated same code in commit E
- Alice and Bob need to study conflict and decide on resolution at pull (time-consuming)
- Possibility of introducing bug on master branch (risky)

loops.cpp (commit C)

```
36
37 // TODO: Code very important loop here ASAP
38
39 ...
40
41
42
43 // TODO: Code other very important loop here ASAP
44
```

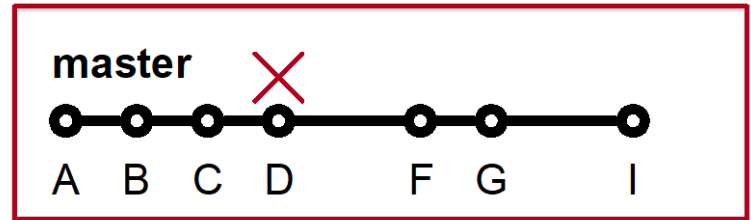
loops.cpp (commit D)

```
36
37 // Very important loop
38 for (int i=0; i<N; ++i) {
39     ...
40     ...
41
42 // Another very important loop
43 for (int i=1; i<=N; ++i) {
44     foo[i] = bar[i] * i;
45 }
```

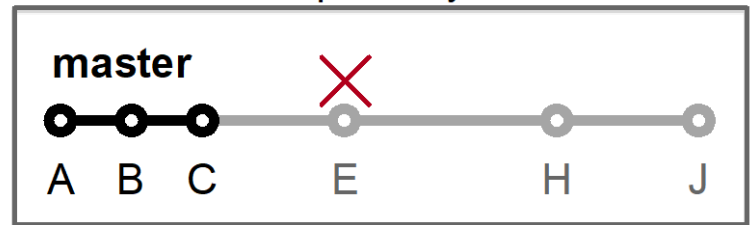
loops.cpp (commit E)

```
36
37 // Very important loop
38 for (int i=0; i<N; i++) {
39     ...
40     ...
41
42 // Another very important loop
43 for (int i=0; i<N; i++) {
44     foo[i] = bar[i] * i;
45 }
```

Alice's Local Repository



Bob's Local Repository



Our First Workflow

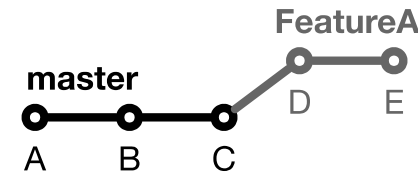
This process of collaborating *via* Git is called the **Centralized Workflow**

- See [Atlassian/BitBucket](#) for more information
- “Simple” to learn and “easy” to use
- Leverages local vs. remote repo dimension
 - Integration in local repo when local repos interact with remote repo
- What if you have many team members?
- What if developers only push once a month?
- What if team members works on different parts of the code?
- Working directly on master

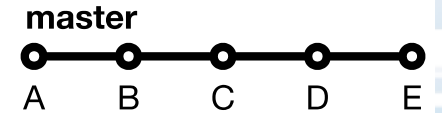
Branches

Branches are independent lines of development

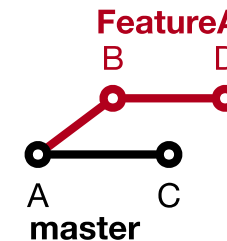
- Use branches to protect master branch
- Feature branches
 - Organize a new feature as a sequence of related commits in a branch
- Branches are usually combined or **merged**
- Develop on a branch, test on the branch, and merge into master
- Integration occurs at merge commits



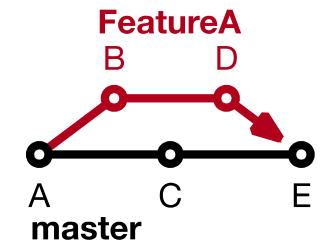
Fast-Forward



No Merge



Divergence

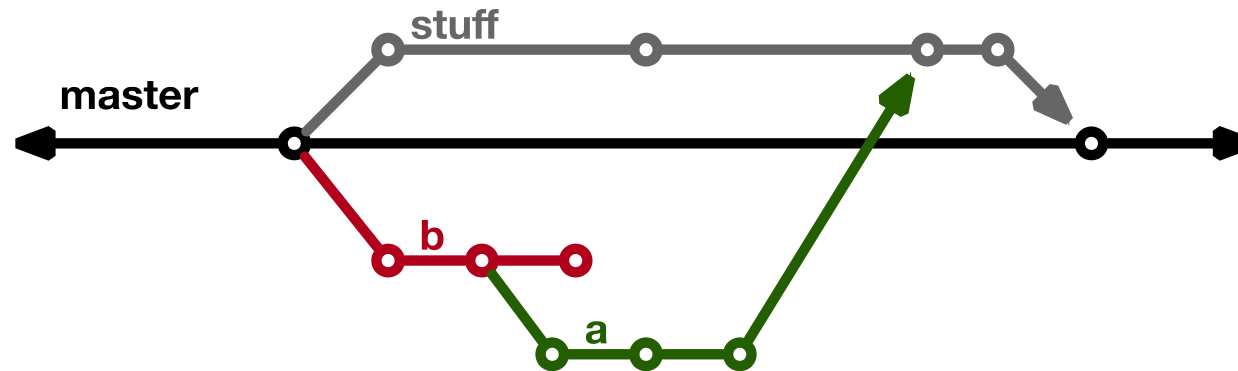


Merge Commit

Control Branch Complexity

Workflow policy is needed

- Descriptive names or linked to issue tracking system
- Where do branches start and end?
- Can multiple people work on one branch?

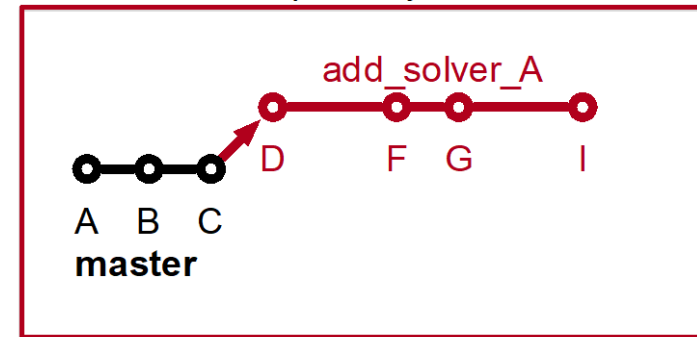


Feature Branches

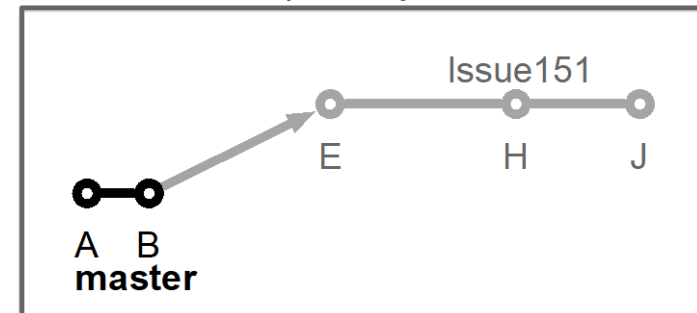
Extend Centralized Workflow

- Remote repo has commits A & B
- Bob pulls remote to synchronize local repo to remote
- Bob creates local feature branch based on commit B
- Commit C pushed to remote repo
- Alice pulls remote to synchronize local repo to remote
- Alice creates local feature branch based on commit C
- Both develop independently on local feature branches

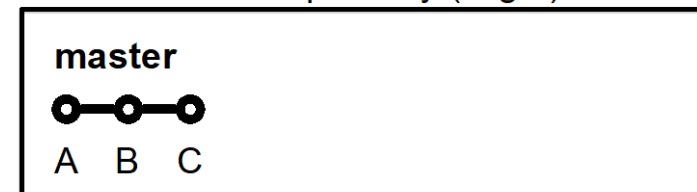
Alice's Local Repository



Bob's Local Repository



Main Remote Repository (origin)

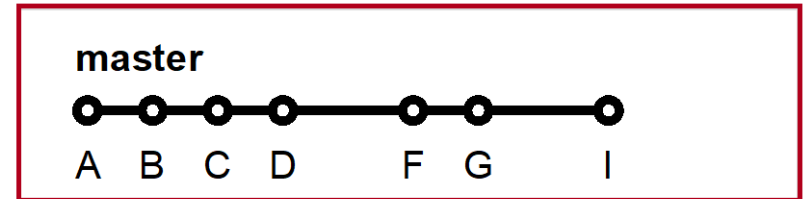


Feature Branch Divergence

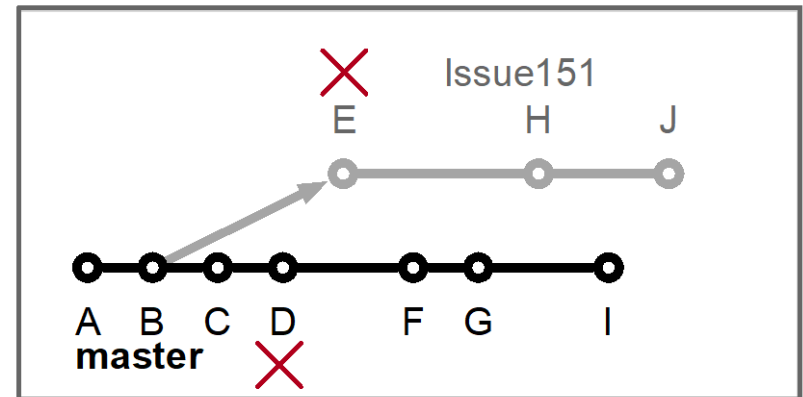
Alice integrates first without issue

- Alice does fast-forward merge to local master
- Alice deletes local feature branch
- Alice pushes master to remote
- Meanwhile, Bob pulls master from remote and finds Alice's changes
- Merge conflict between commits D and E

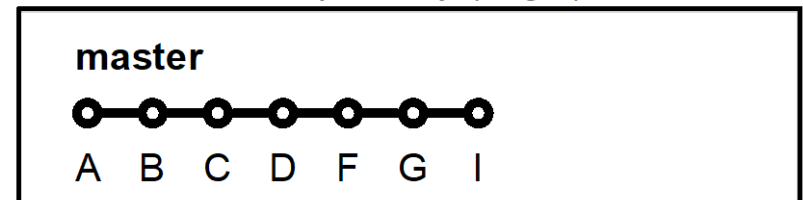
Alice's Local Repository



Bob's Local Repository



Main Remote Repository (origin)

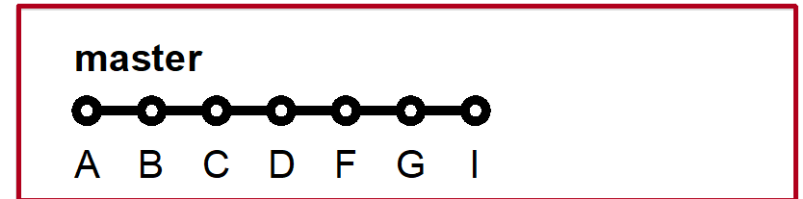


Feature Race Condition

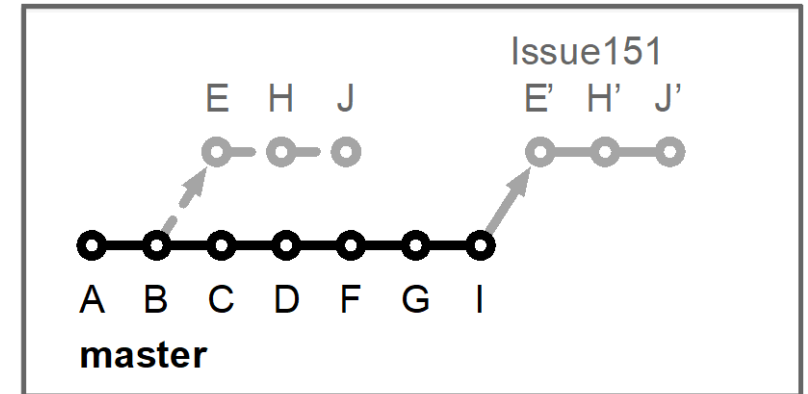
Integration occurs on Bob's local repo

- Bob laments not having fast-forward merge
- Bob **rebases** local feature branch to latest commit on master
 - E based off of commit B
 - E' based off of Alice's commit I
 - E' is E integrated with commits C, D, F, G, I
- Merge conflict resolved by Bob & Alice on Bob's local branch when converting commit E into E'
- Can test on feature branch and merge easily and cleanly

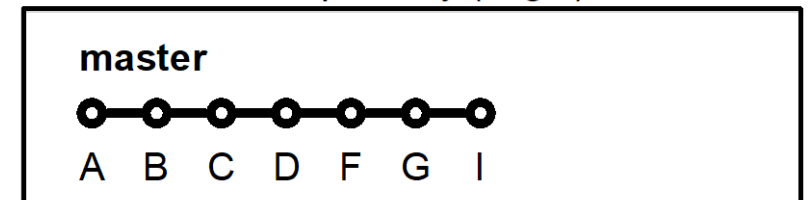
Alice's Local Repository



Bob's Local Repository



Main Remote Repository (origin)



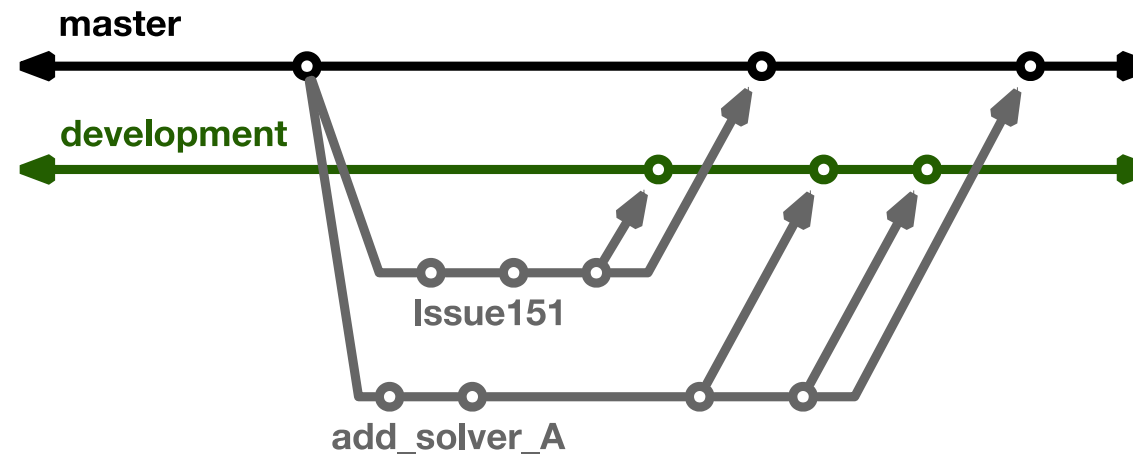
Feature Branches Summary

- Multiple, parallel lines of development possible on single local repo
- Easily maintain local master up-to-date and useable
- Integration with rebase on local repo is safe and can be aborted
- Testing before updating local and remote master branches
- Rebase is advanced Git command
 - Rebase can cause complications and should be used carefully.
- Hide actual workflow
 - History in repo is not represent actual development history
 - Less communication
 - Fewer back-ups using remote repo
- Does it scale with team size? What if team integrates frequently?
- Commits on master can be broken
- See [Atlassian/BitBucket](#) for a richer Feature Branch Workflow

More Branches

Branches with infinite lifetime

- Base off of master branch
- Exist in all copies of a repository
- Each provides a distinct **environment**
 - Development vs. pre-production



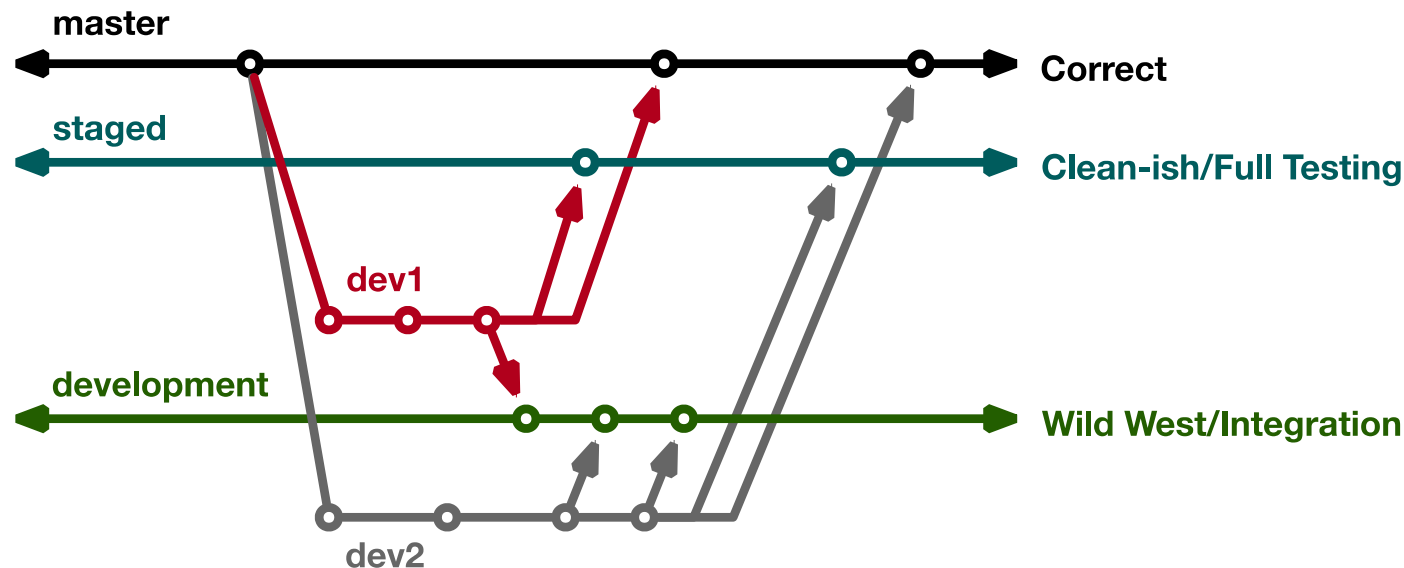
Current FLASH5 Workflow

Test-driven workflow

- Feature branches start and end with master
- All feature branches are merged into development for integration & manual testing
- All feature branches are then merged into staged for full, automated testing

Workflow designed so that

- All commits in master are in staged & development
- infinite branches don't diverge
- Merge conflicts first exposed on development



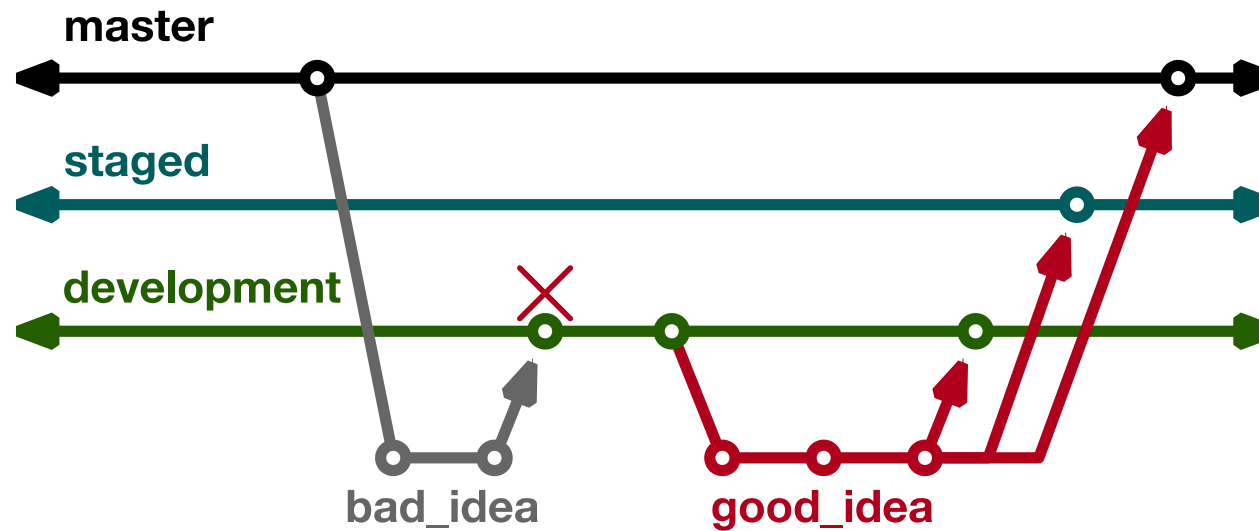
Branch Rules

Why base feature branches off master?

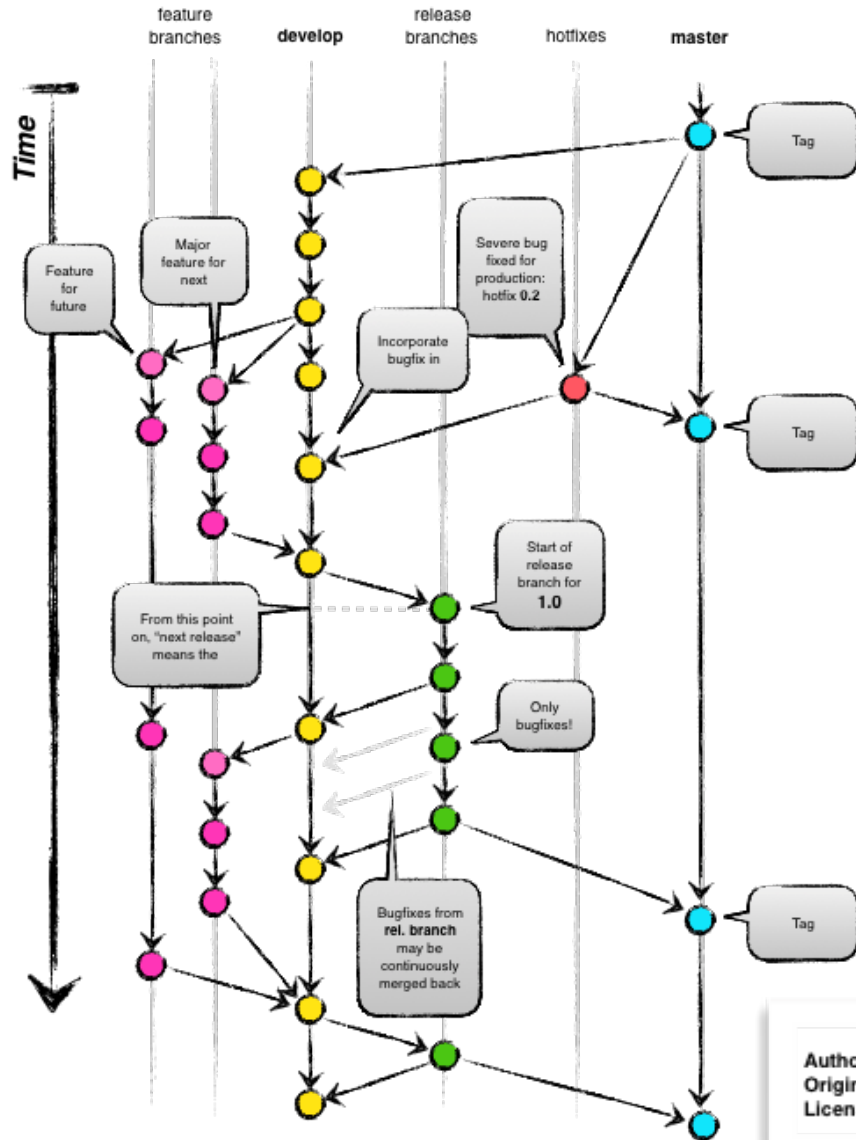
- Start from correct, verified commit
- Clean and simple to learn/enforce
- Isolate master from integration environment

Motivates more rules

- Development never merged into another branch
- Staged never merged into another branch



Git Flow



- Full-featured workflow
- Increased complexity
- Designed for SW with official releases
- Feature branches based off of develop
- Git extensions to enforce policy
- How are develop and master synchronized?
- Where do merge conflicts occur and how are they resolved?

Author: Vincent Driessen
Original blog post: <http://nvie.com/archives/323>
License: Creative Commons



GitHub Flow

<http://scottchacon.com/2011/08/31/github-flow.html>

- Published as viable alternative to Git Flow
- No structured release schedule
- Continuous deployment & continuous integration allows for simpler workflow

Main Ideas

1. All commits in master are **deployable**
2. Base feature branches off of master
3. Push local repository to remote constantly
4. Open Pull Requests early to start dialogue
5. Merge into master after Pull Request review

GitLab Flow

https://docs.gitlab.com/ee/workflow/gitlab_flow.html

- Published as viable alternative to Git Flow & GitHub Flow
- Semi-structured release schedule
- Workflow that simplifies difficulties and common failures in synchronizing infinite lifetime branches

Main Ideas

- Master branch is staging area
- Mature code in master flows downstream into pre-production & production infinite lifetime branches
- Allow for release branches with downstream flow
 - Fixes made upstream & merged into master.
 - Fixes cherry picked into release branch

Things to Think About When Choosing a Git Workflow

Want to establish a clear set of policies that

- results in correct code on a particular branch (usually master),
- ensures that a team can develop in parallel and communicate well,
- minimizes difficulties associated with parallel and distributed work, and
- minimizes overhead associated with learning, following, and enforcing policies.

Adopt what is good for your team

- Consider team culture and project challenges
- Assess what is and isn't feasible/acceptable
- Start with simplest and add complexity where and when necessary

Continuous Integration

The Short & Sweet of Continuous Integration

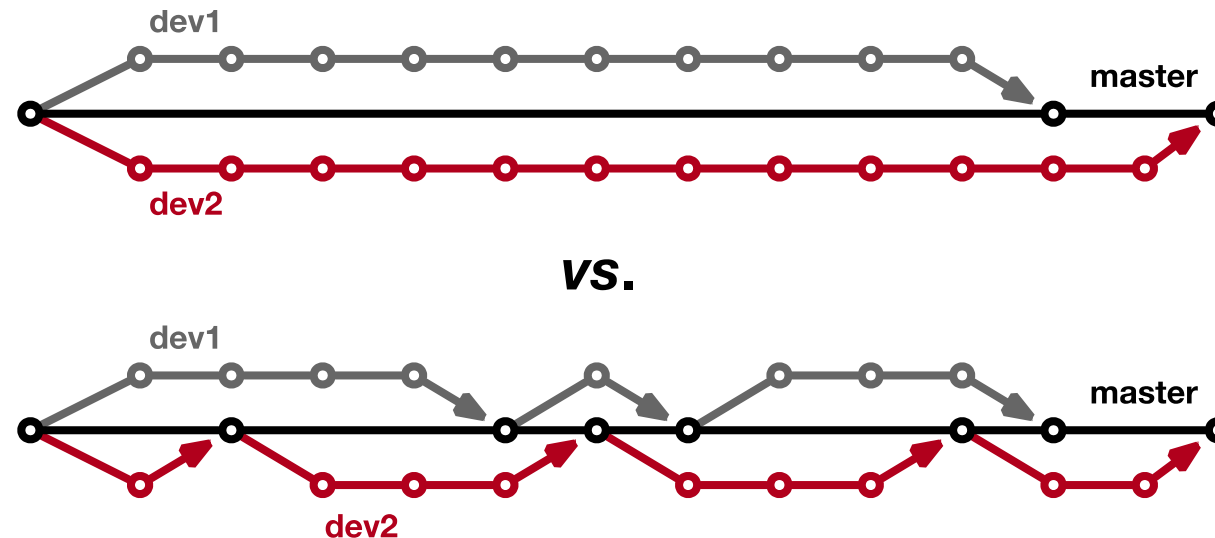
A master branch that always works

- DVCS workflow isolate master from integration environment
- Extend workflow to address difficulties of integrating
 - Minimize likelihood of merge conflict
 - Detect bugs immediately
 - Make debugging process quick and easy

Work Decomposition

Commit and integrate often

- Limit divergence between feature and master branches
- Decreased probability of conflict
- Conflict resolution is simpler and less risky



Error Detection

Test at integration to identify failures immediately

- Control quality of code
- Isolate failure to few commits
- No context switching for programmer

We want a system that

- triggers automated builds/tests on target environments when code changes and
- ideally tests on proposed merge product without finalizing merge.

Test Servers

Servers that

- automate the execution of a test suite or a subset of a test suite,
- allow for running tests on different environments,
- host an interface for viewing results, and
- allows for configuring when the tests are run.

Examples

- CTest/CDash
- Jenkins
- Travis CI and GitLab CI

Cloud-based Test Servers

- Linked to VCS hosts
 - GitHub & Travis CI
 - GitLab CI
 - BitBucket Pipelines
- Automated builds/tests triggered *via* pushes and pull requests
- Builds/tests can be run on cloud systems
- Test results are reported in repository's web interface
- Can trigger code coverage analysis & documentation build

Continuous integration (CI)

- Has existed for some time and interest is growing
- HPC community working to adapt CI for HPC machines
- Setup, maintenance, and monitoring required
- Prerequisites
 - A reasonably automated build system
 - An automated test system with significant test coverage & useful feedback
 - Builds/tests must finish in reasonable amount of time
 - Ability to bundle subset of tests

CI Hello World

Simplest CI example

https://github.com/jrdoneal/CI_HelloWorld

https://travis-ci.org/jrdoneal/CI_HelloWorld

CI example w/ multiple platforms and specific compiler versions

https://github.com/jrdoneal/CI_Multiplatform

Code coverage, testing and CI tutorial (C++)

<https://github.com/amklinv/morpheus>

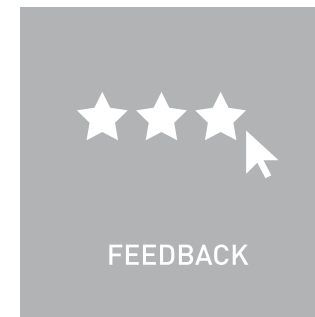
Code coverage, testing, and CI example (Fortran, C++)

<https://github.com/jrdoneal/infrastructure>

Agenda

Time	Module	Topic	Speaker
2:00pm-2:40pm	01	Overview of Best Practices in HPC Software Development	Anshu Dubey, ANL
2:40pm-3:20pm	02	Better (Small) Scientific Software Teams	David E. Bernholdt, ORNL
3:20pm-4:00pm	03	Improving Reproducibility through Better Software Practices	David E. Bernholdt, ORNL
4:00pm-4:30pm		<i>Break</i>	
4:30pm-5:15pm	04	Verification & Refactoring	Anshu Dubey, ANL
5:15pm-6:00pm	05	Git Workflow & Continuous Integration	Jared O'Neal, ANL

<https://r.isc-hpc.com/tut130>
(Please note the R before our domain)



Additional Software-Related Events at ISC 2019

- Tuesday — BOF — **Spack Community BOF**
- Tuesday - Presentation - Parallel Programming - Painful or Productive?
- Tuesday - Poster - (WHPC04) Research Software Engineering enabling HPC
- Tuesday - Poster - (PP06): The International HPC Certification Program
- Tuesday - Poster - (PP09): EPEEC: Productivity at Exascale
- Tuesday — Poster — **(PP24): Helping You Improve Software Sustainability and Development Productivity: An Overview of the IDEAS Project**
- Wednesday — Focus Session — New Approaches, Algorithms Towards Exascale Computing
- Wednesday — BOF — **Software Engineering and Reuse in Computational Science and Engineering** (additional details: <http://bit.ly/swe-cse-bof>)
- Wednesday - BOF - Performance Portability and Productivity: Panel Discussion

Note: **Bold font** denotes events (co-)organized by the IDEAS Productivity project



CI Hello World – Backup Slides

GitHub Repository Page

https://github.com/jrdoneal/CI_HelloWorld

jrdoneal / CI_HelloWorld

Unwatch

1

Star

0

Fork

0

<> Code

Issues 0

Pull requests 0

Projects 0

Wiki

Insights

Settings

No description, website, or topics provided.

Edit

Manage topics

5 commits

1 branch

0 releases

0 contributors

Branch: master

New pull request

Create new file

Upload files

Find file

Clone or download

Developer D. Develop

This change should lead to a correct build environment for the purpos...

Latest commit 93a75c4 2 days ago

.travis.yml	This change should lead to a correct build environment for the purpos...	2 days ago
README.md	Add README file to explain the intent and eventual content of this tu...	2 days ago
hello_world.sh	Add the script that tests that the build environment is correctly con...	2 days ago

Travis CI Configuration File

.travis.yml

```
env:
- TRAVIS_CI_ENV="Hello, World"

#before_install:
#- Put commands here to prepare for executing builds/installs
#- Examples would be using apt-get to install dependencies not
#  included in the Travis CI build environment by default.

#install:
#- Put build commands here
#- In each phase, you can execute multiple commands
#- Travis CI stops if any single command fails in this phase

before_script:
- echo $TRAVIS_CI_ENV

script:
- $TRAVIS_BUILD_DIR/hello_world.sh
#- Travis CI will run each command in this phase even if a previous command
#  terminated in failure

after_success:
- echo "You should see that Hello, World was printed by before_script"

after_failure:
- echo "Hello, World should not have been printed by before_script"
```

The Script Phase

hello_world.sh

```
#!/bin/bash

if [ -z "${TRAVIS_CI_ENV}" ]; then
    echo "Please set the TRAVIS_CI_ENV environment variable"
    exit 1
elif [ "${TRAVIS_CI_ENV}" != "Hello, World" ]; then
    echo "TRAVIS_CI_ENV value is ill-suited for this tutorial"
    exit 2
fi
```

Connecting GitHub & Travis CI

MY ACCOUNT



jrdoneal

Sync account

ORGANIZATIONS

You are not currently a member of any organization.

MISSING AN ORGANIZATION?

Review and add your authorized organizations.



jrdoneal

@jrdoneal

Repositories

Settings

We're only showing your public repositories. You can find your private projects on travis-ci.com.

Legacy Services Integration



Filter repositories



CI_HelloWorld



Settings



CI_Multiplatform



Settings



infrastructure



Settings

Repository in Travis CI

https://travis-ci.org/jrdoneal/CI_HelloWorld

 jrdoneal / CI_HelloWorld  

Current Branches Build History Pull Requests

More options 

✓ **master** This change should lead to a correct build environment for the pu 🔗 #3 passed

tutorial. Travis CI builds should now be successful.

🕒 Ran for 18 sec

📅 27 a day ago

🔗 Commit 93a75c4 

🔗 Compare ff52718...93a75c4 

🔗 Branch master 

 jrdoneal


 </> Ruby

📦 TRAVIS_CI_ENV="Hello, World"

🔄 Restart build


Commit History

.travis.yml
added →


 [jrdoneal](#) / [CI_HelloWorld](#)

[Code](#) [Issues 0](#) [Pull requests 0](#) [Projects 0](#) [Wiki](#) [Insights](#)


Branch: **master** ▼

 Commits on Nov 3, 2018


This change should lead to a correct build environment for the purpos... [...](#)

 Developer D. Develop committed 2 days ago ✓


Update Travis CI configuration file so that it is a step closer to se... [...](#)

 Developer D. Develop committed 2 days ago ✗


Add Travis CI configuration file. With the present content, the build [...](#)

 Developer D. Develop committed 2 days ago ✗

Add the script that tests that the build environment is correctly con... [...](#)

 Developer D. Develop committed 2 days ago

Add README file to explain the intent and eventual content of this tu... [...](#)

 Developer D. Develop committed 2 days ago

Travis CI Build History

Add Travis CI configuration file. With the present content, the build ...


 Developer D. Develop committed 2 days ago ✖

```
▶ 1 Worker information worker_info
▶ 6 Build system information system_info
413
414
415 Setting APT mirror in /etc/apt/sources.list: http://us-east-1.ec2.archive.ubuntu.com/ubuntu/
416
▶ 417 $ git clone --depth=50 --branch=master https://github.com/jrdoneal/CI_HelloWorld.git jrdoneal/CI_HelloWorld git.checkout 0.54s
▶ 427 $ rvm use default rvm 5.27s
▶ 434 $ ruby --version ruby.versions
442 No Gemfile found, skipping bundle install
▼ 443 $ echo $TRAVIS_CI_ENV before_script 0.00s
444
445
446 $ $TRAVIS_BUILD_DIR/hello_world.sh 0.00s
447 Please set the TRAVIS_CI_ENV environment variable
448
449
450 The command "$TRAVIS_BUILD_DIR/hello_world.sh" exited with 1.
▶ 451 $ echo "Hello, World should not have been printed by before_script" after_failure 0.00s
454
455 Done. Your build exited with 1.
```

Top ▲

Travis CI Build History


Update Travis CI configuration file so that it is a step closer to se... ...

 Developer D. Develop committed 2 days ago ✗

```
▶ 1 Worker information worker_info
▶ 6 Build system information system_info
413
414
415 Setting APT mirror in /etc/apt/sources.list: http://us-east-1.ec2.archive.ubuntu.com/ubuntu/
416
▶ 417 $ git clone --depth=50 --branch=master https://github.com/jrdoneal/CI_HelloWorld.git jrdoneal/CI_HelloWorld git.checkout 0.52s
427
428 Setting environment variables from .travis.yml
429 $ export TRAVIS_CI_ENV="This content will result in failure"
430
▶ 431 $ rvm use default rvm 4.53s
▶ 438 $ ruby --version ruby.versions
446 No Gemfile found, skipping bundle install
▼ 447 $ echo $TRAVIS_CI_ENV before_script 0.00s
448 This content will result in failure
449
450 $ $TRAVIS_BUILD_DIR/hello_world.sh 0.00s
451 TRAVIS_CI_ENV value is ill-suited for this tutorial
452
453
454 The command "$TRAVIS_BUILD_DIR/hello_world.sh" exited with 2.
▶ 455 $ echo "Hello, World should not have been printed by before_script" after_failure 0.00s
458
459 Done. Your build exited with 1.
```


Travis CI Build History

This change should lead to a correct build environment for the purpos... ...

 Developer D. Develop committed 2 days ago ✓

```
▶ 1 Worker information worker_info
▶ 6 Build system information system_info
413
414
415 Setting APT mirror in /etc/apt/sources.list: http://us-east-1.ec2.archive.ubuntu.com/ubuntu/
416
▶ 417 $ git clone --depth=50 --branch=master https://github.com/jrdoneal/CI_HelloWorld.git jrdoneal/CI_HelloWorld git.checkout 0.53s
427
428 Setting environment variables from .travis.yml
429 $ export TRAVIS_CI_ENV="Hello, World"
430
▶ 431 $ rvm use default rvm 4.69s
▶ 438 $ ruby --version ruby.versions
446 No Gemfile found, skipping bundle install
▼ 447 $ echo $TRAVIS_CI_ENV before_script 0.00s
448 Hello, World
449
450 $ $TRAVIS_BUILD_DIR/hello_world.sh 0.00s
451
452
453 The command "$TRAVIS_BUILD_DIR/hello_world.sh" exited with 0.
▶ 454 $ echo "You should see that Hello, World was printed by before_script" after_success 0.00s
457
458 Done. Your build exited with 0.
```

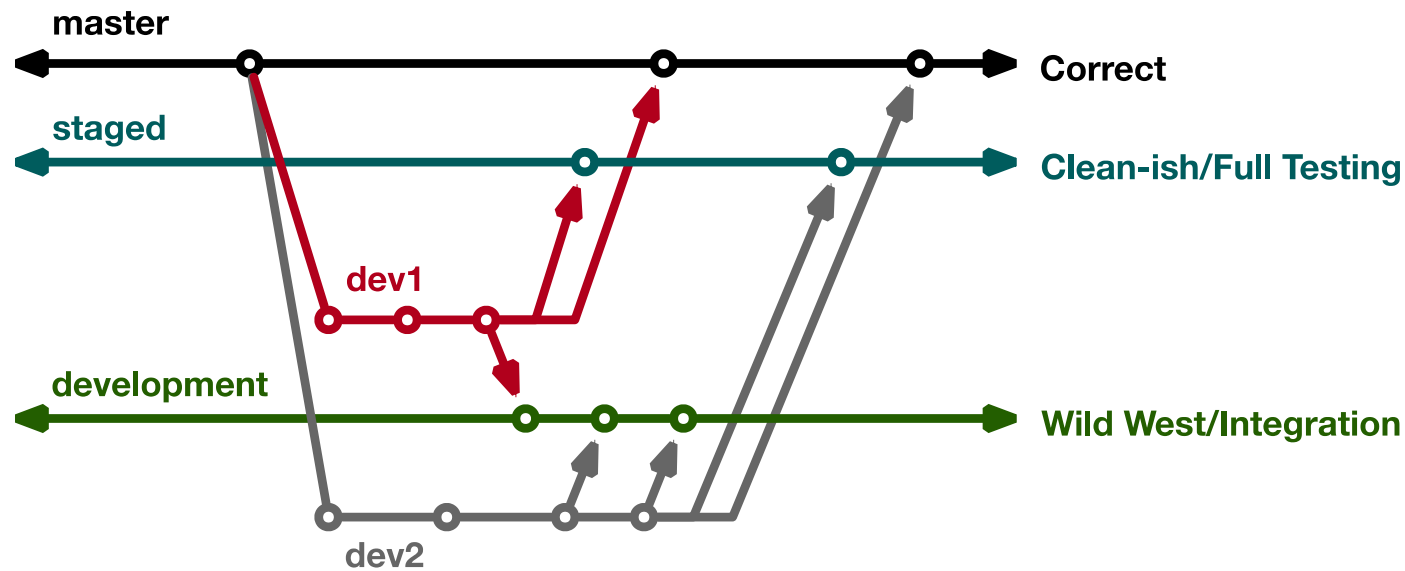
! →

Extra Slides

More Branch Rules

Is staged really necessary?

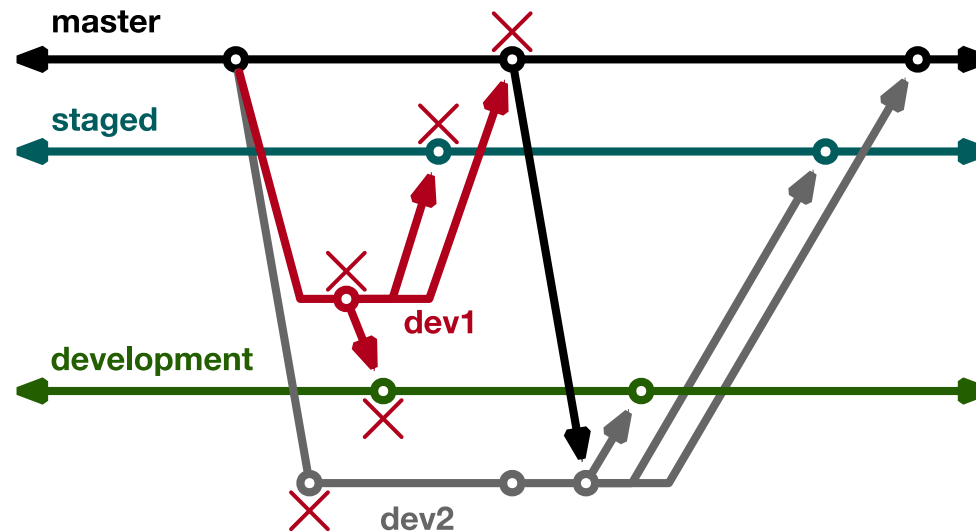
- Contains only changes intended for master
- No integration means cleaner branch
- Allows for extra stage of testing with more tests
- Extra buffer for protecting master branch



Merge Conflicts

How are merge conflicts resolved in FLASH5 Workflow?

- Merge conflict with master means merge conflict with staged and development
- We want to avoid conflict resolution when merging into master
- Directly on feature branch if resolution is there
- One idea is to merge master into feature branch



How do we determine what other tests are needed?

Code coverage tools

- Expose parts of the code that aren't being tested
- gcov
 - standard utility with the GNU compiler collection suite
 - Compile/link with `-coverage` & turn off optimization
 - counts the number of times each statement is executed
- lcov
 - a graphical front-end for gcov
 - available at <http://ltp.sourceforge.net/coverage/lcov.php>
- Hosted servers (e.g. coveralls, codecov)
 - graphical visualization of results
 - push results to server through continuous integration server

Code Coverage Output

Overall Analysis

SOURCE FILES ON BUILD 45					
LIST 2	CHANGED 0	SOURCE CHANGED 0	COVERAGE CHANGED 0		
▲ COVERAGE	Δ	FILE	◆ LINES	◆ RELEVANT	◆ COVERED
— 74.39		src/functions/linear_fcn_class.f90	301	82	61
— 100.0		src/general/modulo_mod.f90	52	3	3

Detailed Analysis

```
265      ! Error distribution same for all x values
266      delta = S*Sxx - Sx*Sx
267      if (delta == 0.0_wp) then
268          ERRORMSG("Cannot do linear least-sqrs. Divide by zero.")
269          stop
270      end if
271      delta_inv = 1.0_wp / delta
```

<https://github.com/jrdoneal/infrastructure>

Code Coverage is Popular

- gcov also works for C and Fortran
- Other tools exist for other languages
 - JCOV for Java
 - Coverage.py for python
 - Devel::Cover for perl
 - profile for MATLAB
 - *etc.*

Special Notes for Morpheus Tutorial

- A code coverage and testing tutorial can be found at the Morpheus repository doxygen pages
 - <https://amklinux.github.io/morpheus/index.html>
- **STEP 1:** These exercises must be run on your own local machine or on a remote machine that you have access to.
- If you cannot generate your own gcov output, the associated lcov output is online
 - <https://amklinux.github.io/morpheus/lcovFiles/index.html>