



Software Packaging



David E. Bernholdt (he/him)
Oak Ridge National Laboratory



Better Scientific Software tutorial @ NOAA Global Systems Laboratory

Contributors: David E. Bernholdt (ORNL), David M. Rogers (ORNL),
IDEAS-ECP Team



See slide 2 for
license details

License, Citation and Acknowledgements

License and Citation

- This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/) (CC BY 4.0).
- **The requested citation the overall tutorial is:** David E. Bernholdt, Anshu Dubey, and Patricia A. Grubel, Better Scientific Software tutorial, in NOAA Global Systems Laboratory, Boulder, Colorado, 2023.
DOI: [10.6084/m9.figshare.23796606](https://doi.org/10.6084/m9.figshare.23796606).
- Individual modules may be cited as *Speaker, Module Title*, in *Tutorial Title*, ...




Acknowledgements

- This work was supported by the U.S. Department of Energy Office of Science, Office of Advanced Scientific Computing Research (ASCR), and by the Exascale Computing Project (17-SC-20-SC), a collaborative effort of the U.S. Department of Energy Office of Science and the National Nuclear Security Administration.
- This work was performed in part at the Argonne National Laboratory, which is managed by UChicago Argonne, LLC for the U.S. Department of Energy under Contract No. DE-AC02-06CH11357.
- This work was performed in part at the Lawrence Livermore National Laboratory, which is managed by Lawrence Livermore National Security, LLC for the U.S. Department of Energy under Contract No. DE-AC52-07NA27344.
- This work was performed in part at the Los Alamos National Laboratory, which is managed by Triad National Security, LLC for the U.S. Department of Energy under Contract No. 89233218CNA000001
- This work was performed in part at the Oak Ridge National Laboratory, which is managed by UT-Battelle, LLC for the U.S. Department of Energy under Contract No. DE-AC05-00OR22725.
- This work was performed in part at Sandia National Laboratories. Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.

Outline

- Why package?
- General Guidelines & Themes
- Simple Walk-Throughs
 - python package
 - Fortran – cmake exports
 - C++ – cmake exports
 - Spack
- Containers
- Performance portability concerns?
- Real-World Examples
 - DCA++: cuda2hip compatibility layer
 - ZFP: scikit-build for cython
 - Cabana: Kokkos with spack

Why Package?



- What does it do?
- How do I set it up?
- Automation *can* be good...
 - but requires really great documentation!

Why Package?

- Standards and conventions save everyone time



1. plug into wall
2. put stuff in top
3. push button
4. take stuff out

Guidelines & Themes

- Start from a portable build system
- Keep source and documentation together
 - So changes are synced
 - YMMV: LAMMPS does this, pyscf does not
- Keep source and tests together
 - Note: some projects maintain separate "reference artifact" repositories
- Split (and separately package) projects that become large
 - Especially true for "optional" components and abstraction layers (aka. "glue-code")

Guidelines & Themes

- Do: Have a CI-level integration test (simulate an external user)
- Do: document manual install process – *what steps do you actually run?*
 - Many projects do this even for dependencies (especially difficult ones)
 - Example: PIConGPU documents how to install Boost (great – since boost has many options)
 - Example: DFT-FE documents how to install Deal.II (great - since Deal.II is complex)
 - Example: lots more inside .github/workflows folders
- Don't: assume everyone will have access to apt-get / docker / VM for getting dependencies
--- as a package consumer ---
- Do: Complain (politely) when something doesn't compile / install / run as documented
 - These are vital fixes and the devs will (should) thank you.
- Do: submit issues / PRs for docs for upstreams
 - Great way to make friends & forge collaborations.

Package Publication Checklist

pre-flight checks

- Is this something I am going to re-use?
- Is the documentation good enough that another developer can quickly get it working?
- Can I hold development of new features while I package up what's here?
 - "pausing" a good idea is nontrivial
- Have I tested it in practice? – start from a clean copy, follow the directions / tests
- Am I ready to support users of this software? (or write a disclaimer)
- Have I picked a license and figured out what copyright assignment & internal reviews need to happen.
- Have I documented my git workflow (what do branches / tags represent)?

Simple Walk-Throughs

- Python - pyscaffold
- Fortran – CMake Library Export
- C++ - CMake Library Export
- cmake with spack

Hello Numerical World Example (heat equation)

github.com/bssw-tutorial/simple-heateq

```
33 COPYING
56 README
36 CMakeLists.txt
20 Makefile
13 build.sh
--> src/
    143 src/pheat.py
    192 src/cheat.cc
    269 src/fheat.f90
```

- Minimal working code for each language: parameter class, energy/integrator class, and main function
- Time to build up the developer and user interfaces!

Hello Numerical World Example (heat equation)

- How will other projects use this work?

```
33 COPYING
56 README
36 CMakeLists.txt
20 Makefile
13 build.sh
--> src/
    143 src/pheat.py
    192 src/cheat.cc
    269 src/fheat.f90
```

Front-lines: Documentation!

- * what's expected to work?
- * where / how do I configure it?

executable

```
$PREFIX/bin/
    artifact-tools
    run-parallel
    run-serial
```

headers

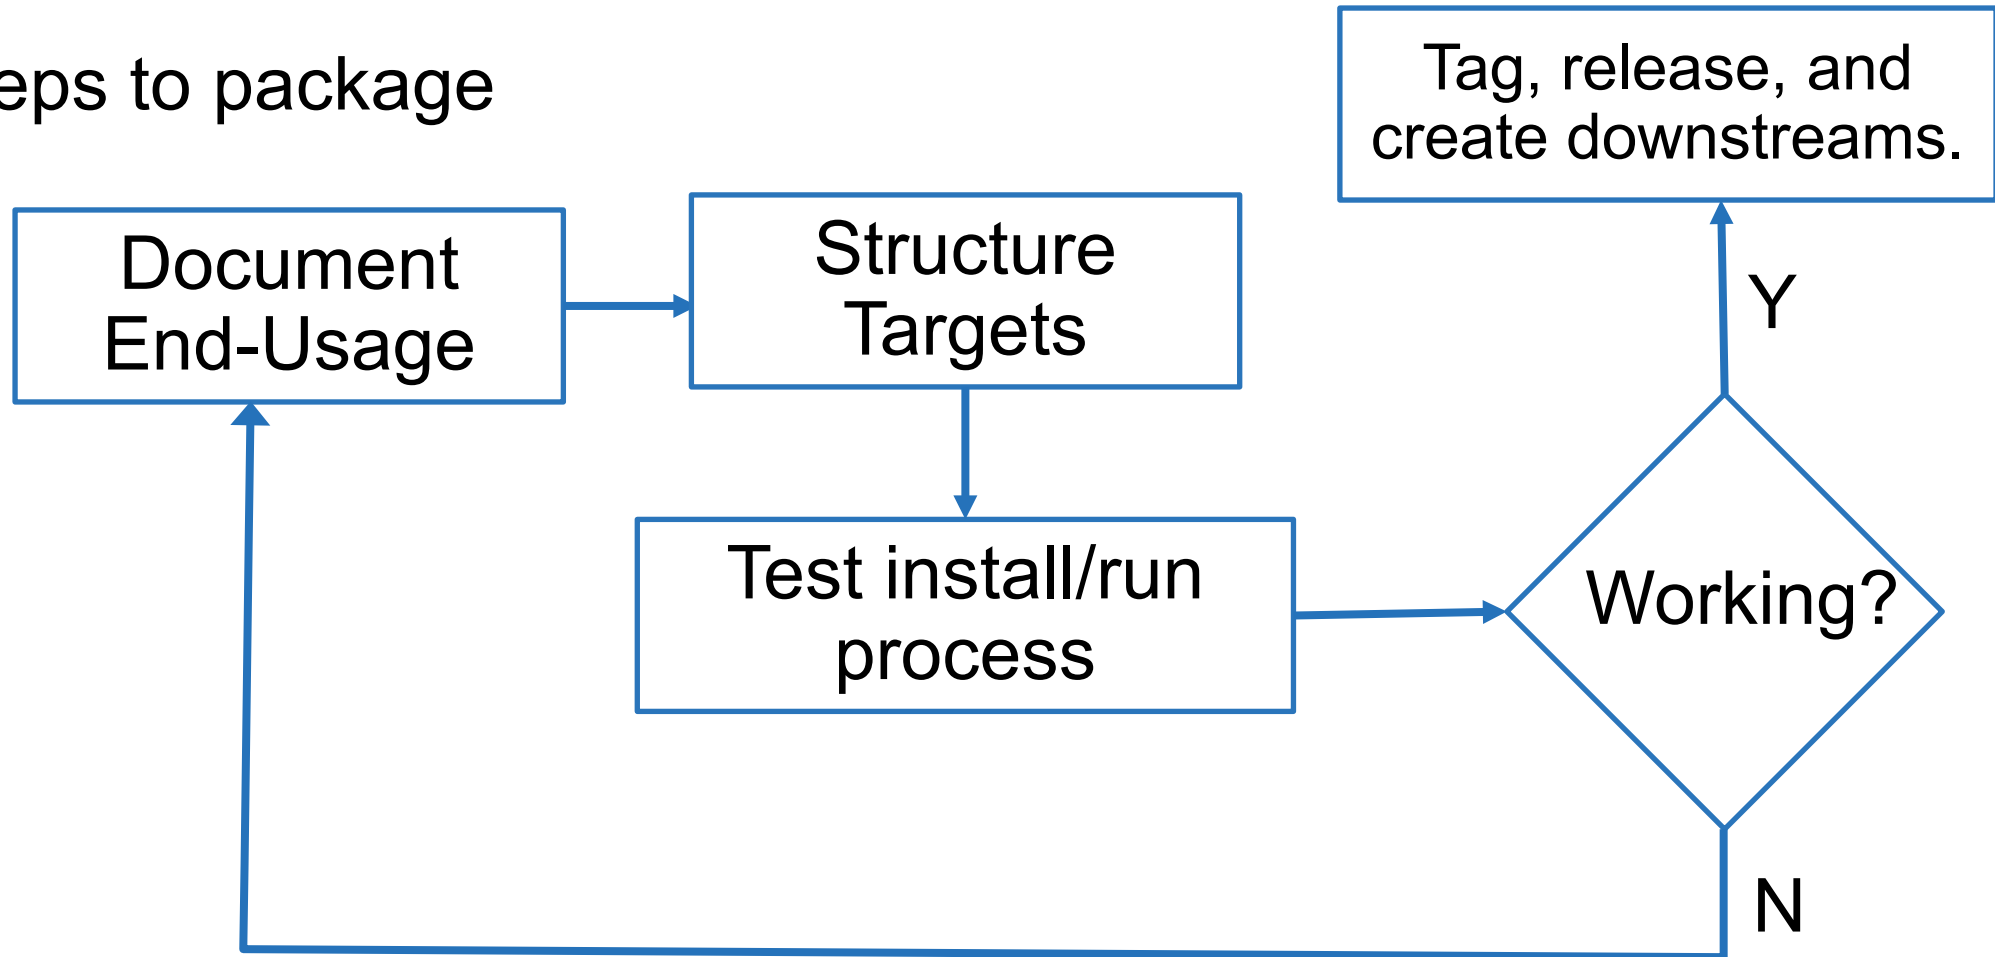
```
$PREFIX/include/$PROJ
    config.h
    heat.h
    heat.mod
```

libraries

```
$PREFIX/lib/$PROJ
    libheat.so
    libheat.a
```

Hello Numerical World Example (heat equation)

- Steps to package



Importing a Python Package

basic

```
# requirements.txt
heateq >= 0.1
```

```
$ pip install -r requirements.txt
# or
$ export PYTHONPATH=/path/to/heateq
$ python3 app.py
```

advanced

```
# setup.cfg
```

```
install_requires =
    heateq >= 0.1
```

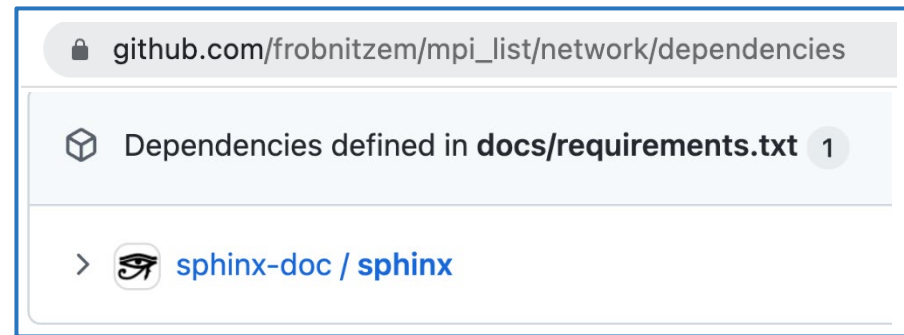
```
$ python -m venv venv
$ source venv/bin/activate
(venv) $ pip install -e .
(venv) $ python3
>>> import app
>>>
```

```
# app.py
```

```
import heateq
```

```
# app.py
```

```
from heateq.heat import Params
```



Python Library Structure

<prefix>/lib/python3.x/site-packages/

- src/pheat.py class Params
 class Energy -copy--> heateq/pheat.py
 def simulate(p)
- __init__.py (can be empty) -copy--> heateq/__init__.py

|
v

Inside the heateq package:
from .pheat import simulate

Outside the package:
from heateq.pheat import simulate

Packaging with pyscaffold

```
$ pip3 install pyscaffold
$ pip3 install tox
$ putup heateq
$ cd heateq # tests in tests/ subdir.
$ tox
```

```
default run-test: commands[0] | pytest
===== test session starts =====
platform darwin -- Python 3.9.0, pytest-6.2.2, py-1.10.0, pluggy-0.13.1 -- plugins:
cov-2.11.1
collected 2 items
```

```
tests/test_skeleton.py::test_fib PASSED [ 50%]
tests/test_skeleton.py::test_main PASSED [100%]
```


```
----- coverage: platform darwin, python 3.9.0-final-0 -----
Name           Stmts  Miss Branch BrPart  Cover  Missing
```

```
-----
src/heatq/__init__.py      6     0     0     0  100%
src/heatq/skeleton.py    32     1     2     0   97%  135
-----
```

```
TOTAL                38     1     2     0   98%
```

```
===== 2 passed in 0.07s =====
```

```
default: commands succeeded
congratulations :)
```



pyscaffold.org

Net result

```
33 COPYING
56 README
36 CMakeLists.txt
20 Makefile
13 build.sh
--> src/
    143 src/pheat.py
    192 src/cheat.cc
    269 src/fheat.f90
```

```
33 COPYING.rst
80 README.rst
 5 AUTHORS.rst
13 CHANGELOG.rst
 8 pyproject.toml
68 tox.ini
21 setup.py
100 setup.cfg
    docs/
    tests/
36 CMakeLists.txt
20 Makefile
13 build.sh
--> src/
    143 heateq/pheat.py
    192 cheat.cc
    269 fheat.f90
```

- **setup.cfg**: editable list of project data & dependencies
- **pyproject.toml**, **tox.ini**, **setup.py**: auto-generated boilerplate
- README: note "pip -e install ." command

Importing a Fortran Package

basic

```
gfortran -I$inst/include/heateq \
  -L$inst/lib \
  -Wl,-rpath,$inst/lib -lfheateq \
  -o app app.f90
```

```
/* app.f90 */
program app
  use ArgParser
  use EnergyField
```

...

advanced

```
# CMakeLists.txt
option(ENABLE_HEATEQ "Use heateq library." ON)

if(ENABLE_HEATEQ)
  find_package(heateq 1.0 REQUIRED)
  target_link_libraries(app PRIVATE
    heateq::fheateq)
endif()
```

```
/* app.hpp.in */

#cmakedefine ENABLE_HEATEQ
```

Fortran Library Structure

- `src/fheat.f90`
 - `gfortran -shared --->`
 - `module ArgParser` -----> `include/heateq/argparser.mod`
 - `module EnergyField` -----> `include/heateq/energyfield.mod`
 - `use ArgParser`
 - > `lib/libfheateq.a`

|
v

Requires referencing correctly

`use EnergyField`

```
gfortran -I$inst/include/heateq \
-L$inst/lib \
-Wl,-rpath,$inst/lib -lfheateq \
-o app app.f90
```

Complications: Transitive Build / Link Requirements

- Header include paths
- Library search paths
- Compiler features
 - e.g. C++11/14/17/20
 - Compiler-dependent runtimes (GCC OpenMP vs. Clang)
- Linking features
 - Fat-binary formatted coprocessor objects.

LibXYZ

OpenPMD

CUDA

openblas

Heat

Multiphysics

Intended to be solved by (pick one)

- pkgconfig/\$PROJ.pc
- cmake/\$PROJ/\${PROJ}Config.cmake

Installing a Fortran library with CMake

```
# CMakeLists.txt
...
add_library(fheateq
  src/argparser.f90
  src/energyfield.f90)
install(TARGETS fheateq
  DESTINATION lib
  EXPORT heateqTargets)
install(EXPORT heateqTargets
  FILE heateqTargets.cmake
  NAMESPACE heateq::
  DESTINATION lib/cmake/heateq
)
... # 15 more lines of cmake cruft
```

```
# heateq.cmake.in

@PACKAGE_INIT@

include (
  "${CMAKE_CURRENT_LIST_DIR}/heateqTargets.cmake" )

include(CMakeFindDependencyMacro)
find_dependency(MPI 2.0 REQUIRED)

check_required_components(heateq)
```

Installs to <prefix>/lib/cmake/heateq/heateq.cmake

- References:

- github.com/bssw-tutorial/simple-heateq (pkg branch)
- Note: the cruft is done for you in "cmake/install.cmake: install_libs()"

Net result

```
33 COPYING
56 README
36 CMakeLists.txt
20 Makefile
13 build.sh
--> src/
    143  src/pheat.py
    192  src/cheat.cc
    269  src/fheat.f90
```

- fheateq.cmake/pc.in – see examples
- README: note how downstreams should use "find_package(<package name>)" and target_link_libraries()
- example -- simple downstream consumer of this library
- ChangeLog: document your success!

```
33 COPYING
86 README
29 ChangeLog
50 CMakeLists.txt
8  fheateq.cmake.in
9  fheateq.pc.in
--> cmake/
    65  install.cmake
    23  rpath.cmake
20 Makefile
--> example/
    9  CMakeLists.txt
    18  test.sh
13 build.sh
--> src/
    143  pheat.py
    192  cheat.cc
    125  fheat.f90
    44  ArgParser.f90
    98  EnergyField.f90
```

Importing a C++ Package

basic

```
g++ -I$inst/include \
    -L$inst/lib \
    -Wl,-rpath,$inst/lib -lcheateq \
    -o app app.cpp
```

```
/* app.cpp */

#include <heateq/heat.hpp>

...
```

advanced

```
# CMakeLists.txt
option(ENABLE_HEATEQ "Use heateq library." ON)

if(ENABLE_HEATEQ)
    find_package(heateq 1.0 REQUIRED)
    target_link_libraries(app PRIVATE heateq::cheateq)
endif()
```

```
/* app.hpp.in */

#cmakedefine ENABLE_HEATEQ
```

C++ Library Structure

- include/heateq/heat.hpp

```
struct Params {}  
struct Energy {}
```

- src/cheat.cpp

```
struct Params {}  
struct Energy {}
```

----(copy)-----> include/heateq/heat.hpp

--(g++ -shared)--> lib/libcheateq.so

|
|
|
v

#include <heateq/heat.hpp>

Requires referencing correctly

```
g++ -I$inst/include \  
    -L$inst/lib \  
    -Wl,-rpath,$inst/lib -lcheateq \  
    -o app app.cpp
```

Installing a C++ library with CMake

```
# CMakeLists.txt
...
add_library(cheateq src/cheateq.cpp)
install(TARGETS cheateq
        DESTINATION lib
        EXPORT cheateqTargets)
install(EXPORT cheateqTargets
        FILE cheateqTargets.cmake
        NAMESPACE cheateq::
        DESTINATION lib/cmake/cheateq
)
... # 15 more lines of cmake cruft
```

```
# cheateq.cmake.in

@PACKAGE_INIT@

include (
    "${CMAKE_CURRENT_LIST_DIR}/cheateqTargets.cmake" )

include(CMakeFindDependencyMacro)
find_dependency(MPI 2.0 REQUIRED)

check_required_components(cheateq)
```

- References:

- github.com/frobnitzem/lib0
- <https://code.ornl.gov/99R/mpi-test>
- <https://cmake.org/cmake/help/git-stage/manual/cmake-packages.7.html#creating-packages>
- Can also create a cheateq.pc.in for pkg-config (see *pkg* branch)

Package Publication Steps – C++/Fortran with cmake

- After editing CMakeLists.txt:
- Check and run tests with cmake; make && ctest
- Update ChangeLog, documentation
- `git tag -m "Bug fixes to v1.1.1" v1.1.2`
- `git push`
- Change public facing websites, modules, spack versions, links, etc. to point to new version
 - For spack <package name>/package.py, use "spack checksum <package name>"
- **** Users should find and use new versions ****
 - This highlights the need for testing deployments using both simultaneous versions and update-in-place strategies. Did you document that?

Releases 12

 **libzmq 4.3.4** Latest
on Jan 17, 2021

+ 11 releases

Going further – additional notes and resources

- We've generated test scripts – add them to CI
 - This lets you easily try multiple different build environments (OS, compiler, etc.)
 - You will notice quickly when a new (upstream) dependency release breaks your code.
 - This will naturally lead you to put dependency compatibility versions in build files and docs.
- Some good examples for Fortran package structure:
 - Well documented: <https://github.com/leonfoks/coretran>
 - Namespace conventions: <https://selalib.github.io/>
 - Fortran standard lib (using fypp meta-programming): <https://github.com/fortran-lang/stdlib>
 - Fortran Package Index: <https://fortran-lang.org/>, <https://www.archaeologic.com/software>
 - Fortran Package Manager: <https://fpm.fortran-lang.org/>
 - Alternative / complementary approach to cmake that works well within the Fortran ecosystem

Package Publication Steps – C++ with cmake +



spack.readthedocs.io

- Spack replaces "build.sh" with a spec

```
33 COPYING
84 README
29 ChangeLog
50 CMakeLists.txt
20 Makefile
13 build.sh
--> tests/
    30 test_heat.sh
--> src/
    143 pheat.py
    192 cheat.cc
    269 fheat.f90
```

```
# heateq/package.py

from spack import *

class HeatEq(CMakePackage):
    "HeatEq: heat conduction kernels"
    homepage = "https://..."
    maintainers = ["github-id"]
    def cmake_args(self):
        mpi = self.spec["mpi"]
        return [ "-DMPI_HOME={0}"
                .format(mpi.prefix) ]
```

- README: now references "spack install heateq"
- Eventually: package.py knows how to compile your package's variants and historical versions



Anatomy of a Spack Dependency "spec"

```
193     - e4s_22.02_gpu_specs:
194       # Minimal diff from v21.11
195       - amrex@22.02 +rocm~cuda amdgpu_target=gfx90a
196       - kokkos@3.5.00 +rocm~cuda~wrapper~openmp amdgpu_target=gfx90a
197       - strumpack@6.3.0 ~slate+rocm~cuda amdgpu_target=gfx90a
198       - sundials@6.1.1 +rocm~cuda amdgpu_target=gfx90a
```

<https://github.com/mpbelhorn/olcf-spack-environments/blob/develop/hosts/frontier/envs/base/spack.yaml>

<package name>@<version>

+<enabled option> ~<disabled option>

% <compiler>@<compiler version>

^<dependency1> ^<dependency2> ...

https://spack.readthedocs.io/en/latest/packaging_guide.html#dependency-specs

"Progression" of Packaging

- Build System
 - Automake / scons / cmake / mesonbuild.com
- Package Management
 - Pkg-config / CMake Package Manager / spack
- Containerization
 - Singularity / charliecloud + docker-compose
- References
 - <https://supercontainers.github.io/sc20-tutorial/>
 - https://fluid-run.readthedocs.io/en/latest/HowTo/setup_your_repo.html

Containerization

Xen Hypervisor = kernel built to manage kernels

Linux Kernel

Daemons

User Programs

Real Filesystems

Virtual Machine

- Kernel, Daemons
- User Programs + tty/gui
- Disk Image Filesystem

App Container

- Emulated / shared filesystems + images
- User program(s)

...

FreeBSD
Kernel

Linux
Kernel

...

Virtualization vs. Containerization

Virtual Machines [VirtualBox, KVM+QEMU, ...]

- Act at the OS-level, run their own kernel
- Disk image filesystem (lots of space)
- Some support processor emulation
- Must be self-contained (think network-level connectivity like NFS-mounts)

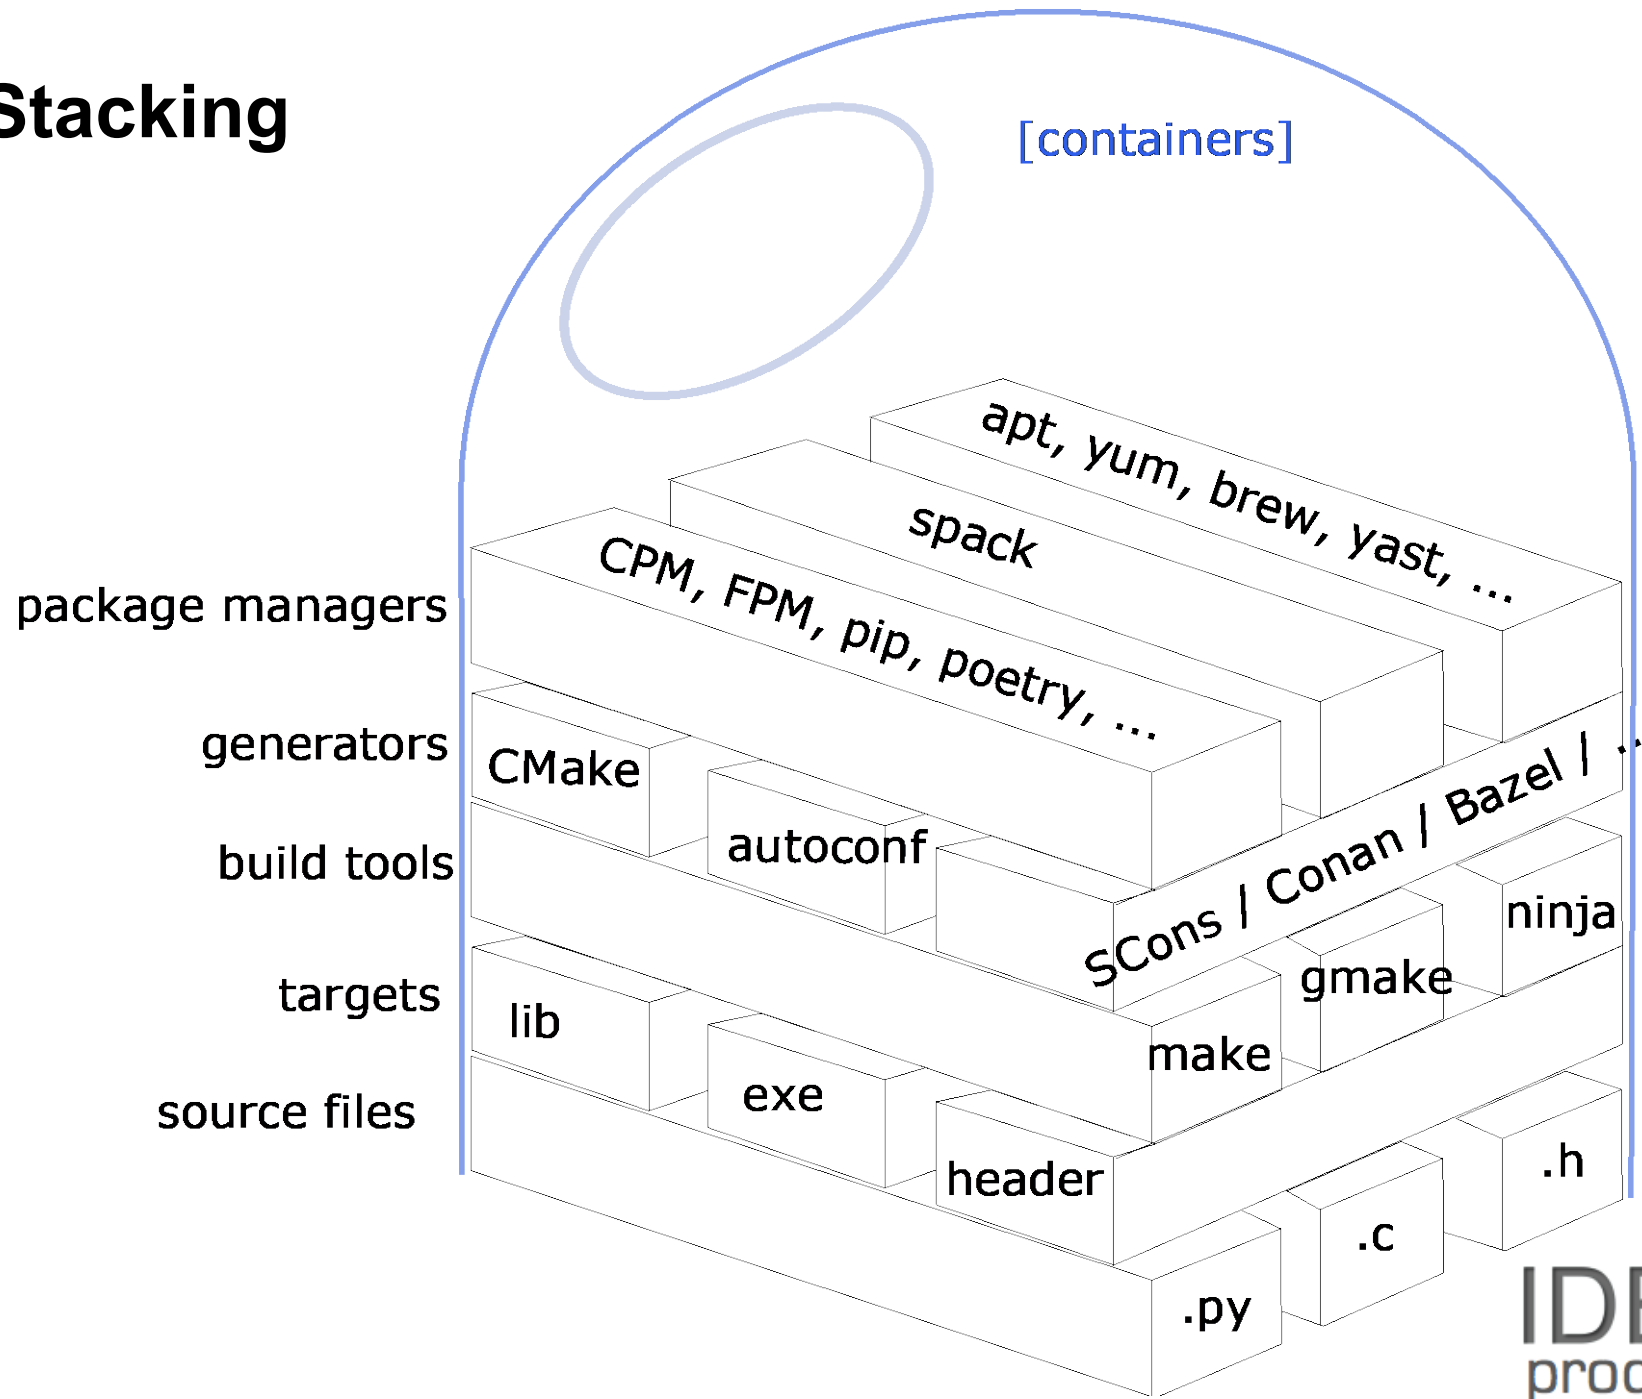
Both:

allow checkpoint / restart

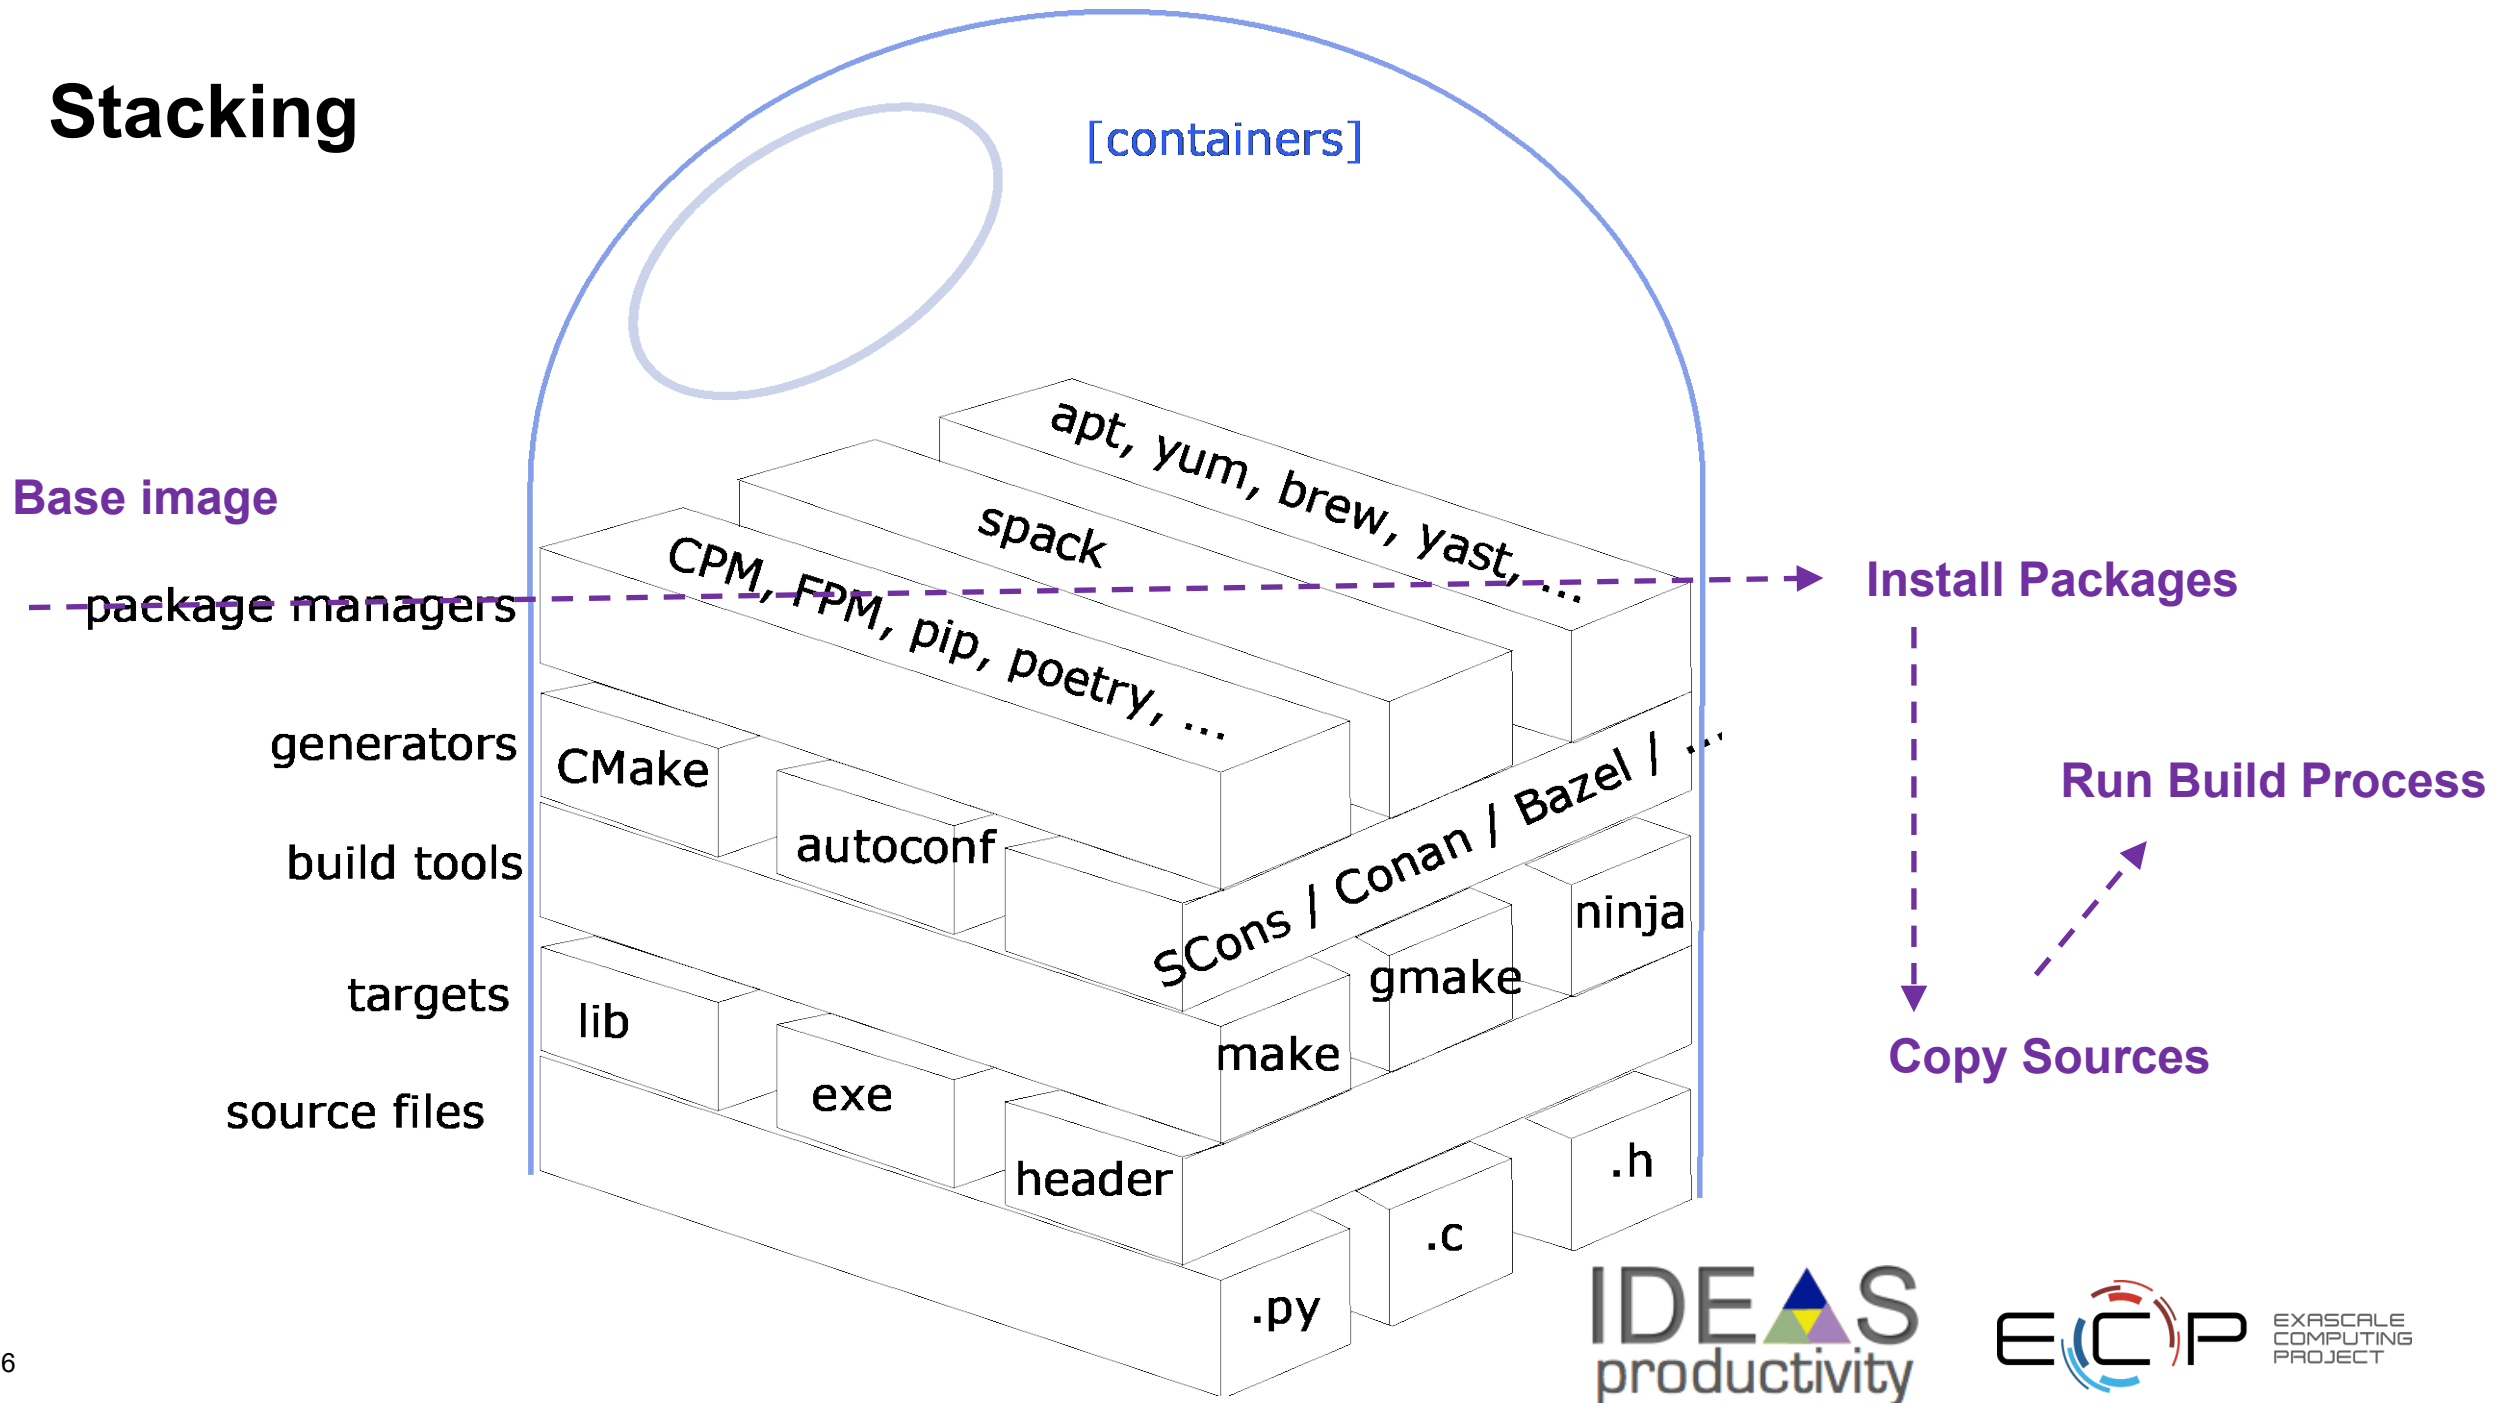
Containers [Docker, Apptainer, Charlie-Cloud, ...]

- Act at the application-level, and share the same OS
- Virtual filesystems = fully custom system libraries, SW stack, and tools
- Can still mount/map libraries and system facilities from host
- Distinguish "image" (stored container) from "container" (running container)

Stacking



Stacking



Conclusion

- Documentation is the beginning and end of packaging
 - Makefiles, dependency lists, and scripts are no substitute for explanations
- Lots of standards & tools to choose from!
 - Make / CMake / autotools
 - py-scaffold / poetry
 - setup.py/"make-ext", scikit-build+cython
 - spack
- Packaging helps you...
 - Interact with your users
 - Improve your developing experience (lower cognitive load)
 - More easily test
 - Deploy faster

Acknowledgments

- IDEAS-ECP Team:
 - David Bernholdt
 - Patricia Grubel
 - Mark Miller
 - Axel Huebl
- PIConGPU Team:
 - Sunita Chandrasekaran
 - Rene Widera
 - Klaus Steiniger
 - Alexander Debus
- DFT-FE Team:
 - Vikram Gavini
 - Sambit Das
 - Phani Motamarri
- DCA++ Team:
 - Peter Doak
 - Thomas Maier
- ZFP Team:
 - Peter Lindstrom
- OLCF/HPE/Spack Teams:
 - Matt Belhorn
 - Luke Roskop
 - Massimiliano Culpo
 - Todd Gamblin

New article on CI team practices:

https://bssw.io/blog_posts/bright-spots-team-experiences-implementing-continuous-integration



Bonus Material

Anatomy of a Spack Dependency "spec"

<package name>@<version>

+<enabled option> ~<disabled option>

% <compiler>@<compiler version>

^<dependency1> ^<dependency2> ...

- compile-time options to your package:
 - `variant('option-name', default=False, description='help text')`
- dependency for your package:
 - `depends_on("spec string", when="string-to-test-against-my-spec") #, type='build'`
- package idiom:

```
def cmake_args(self):  
    spec = self.spec  
    if spec.satisfies("+myoption"): ...
```

Advanced Examples:
spack edit gcc

https://spack.readthedocs.io/en/latest/packaging_guide.html#dependencies-specs

Package Publication Steps – C++ with cmake +



spack.readthedocs.io

```
# heateq/package.py

from spack import *


class HeatEq(CMakePackage):
    "HeatEq: heat conduction kernels"
    homepage = "https://..."

    variant('openmp', default=True)
    depends_on("py-pybind11@2.6.2")
    depends_on('llvm-openmp', when='%apple-clang +openmp')

    def cmake_args(self):
        spec = self.spec
        args = [ "-DMY_BUILD_TESTS=YES"
                "-DENABLE_OPENMP=%s" % (" +openmp" in spec) ]
        return args
```

State of Practice – Packaging for Portability

portability (kernels)



Kokkos

Raja

Alpaka

DPC++ / SYCL

... Thrust

std::execution::par.unseq

Real-World Example: DCA++

- Dynamic Cluster Approximation
 - Electron correlation involving many tensor contractions (matrix multiplies)
 - C++ code
 - Implements own matrix math library, adding HIP backend
- Challenge
 - Minimal additions to existing CUDA build method
 - Several types of link helpers – runtime, blas, kernel
- Solution
 - Use cmake to include a header-translation layer and change link options – minimal changes to source code.

Real-World Example: DCA++

```
// src/linalg/util/info_gpu.cpp

// This file implements gpu info functions.

#include "dca/config/haves_defines.hpp"
#if defined(DCA_HAVE_CUDA)
#include "dca/linalg/util/error_cuda.hpp"
#elif defined(DCA_HAVE_HIP)
#include "dca/linalg/util/error_hip.hpp"
#include "dca/util/cuda2hip.h"
#endif
```

- References:

- <https://github.com/CompFUSE/DCA>
- <https://github.com/twhite-cray/quip>
- https://code.ornl.gov/99R/mpi-test/-/tree/gpu_support

Real-World Example: pyscf extension template

- Python Atomic Orbital Code – HF, DFT, some CC
 - Modular python design
 - Kernels implemented in C for efficiency
 - Extended functionality as plugins (e.g. analysis helpers, MPI parallelization)
- Challenge
 - Enable pyscf to "import" its plugins
 - Allow plugins to incorporate compiled C libraries
- Solution
 - Standardize package layout and provide a templated "setup.py" file.
- References:
 - <https://github.com/pyscf/extension-template>

Real-World Example: pyscf extension template

```
# setup.py
...
def make_ext(pkg_name, srcs,
             libraries=[], library_dirs=[pyscf_lib_dir],
             include_dirs=[], extra_compile_flags=[],
             extra_link_flags=[], **kwargs):
    return Extension(pkg_name, srcs,
                    libraries = libraries,
                    library_dirs = library_dirs,
                    include_dirs = include_dirs + library_dirs,
                    extra_compile_args = extra_compile_flags,
                    extra_link_args = extra_link_flags,
                    runtime_library_dirs = runtime_library_dirs, **kwargs)

if 'SO_EXTENSIONS' in metadata:
    settings['ext_modules'] = [make_ext(k, v) for k, v in SO_EXTENSIONS.items()]
```

- References:

- <https://github.com/pyscf/extension-template>

Real-World Example: ZFP

- Scientific Data Compression Library
 - C++ code
 - Focus is on multidimensional arrays
- Challenge
 - Export all functionality to python with minimal effort
 - C++ code contains non-trivial data structures and link dependencies
- Solution
 - Adopt scikit-build process using cython C++ wrappers
- References:
 - <https://github.com/LLNL/zfp>
 - <https://scikit-build.readthedocs.io>

Real-World Example: ZFP

```
# python/zfpy.pyx
...
cdef bytes compress_numpy(
    np.ndarray arr,
    double tolerance = -1,
    double rate = -1,
    int precision = -1,
    write_header=True
):
    ...

# Setup zfp structs to begin compression
cdef zfp_field* field =
    _init_field(arr)
cdef zfp_stream* stream =
    zfp_stream_open(NULL)
```

```
# python/CMakeLists.txt

...
add_cython_target(zfpy zfpy.pyx C)
```

```
# python/zfpy.pxd

import cython
cimport libc.stdint as stdint

cdef extern from "bitstream.h":
    cdef struct bitstream:
        pass
    bitstream* stream_open(void* data, size_t)
    void stream_close(bitstream* stream)

...
```

• References:

- <https://github.com/LLNL/zfp>
- <https://scikit-build.readthedocs.io>

Real-World Example: Cabana

- <https://github.com/ECP-copa/Cabana>

- Molecular Dynamics (Particle) simulation library
 - C++ code using Kokkos performance portability library
 - Focus is on flexible data layouts for particles
- Challenge
 - Provide a spack compile recipe correctly targeting Kokkos library
 - Allow user-selection of kokkos backends and features to be visible from library
 - Connect to library consumers (MD applications)
- Solution
 - Careful documentation of spack options required from its Kokkos dependency

Real-World Example: Cabana

```
# spack edit cabana
```

```
from spack.pkg.builtin.kokkos import Kokkos
```

```
...
```

```
_versions = {  
    ":0.2.0": "-legacy",  
    "0.3.0": "@3.1:",  
    "0.4.0": "@3.2:"  
}
```

```
for _version, _kk_version in _versions.items():  
    for _backend in kokkos_backends:  
        if (_kk_version == "-legacy" and _backend == 'pthread'):  
            _kk_spec = 'kokkos-legacy+threads'  
        elif (_kk_version == "-legacy" and  
             _backend not in ['serial', 'openmp', 'cuda']):  
            continue  
        else:  
            _kk_spec = 'kokkos{0}+{1}'.format(_kk_version, _backend)  
            depends_on(_kk_spec, when='@{0}+{1}'.format(_version, _backend))
```

HPC: modules and Spack Development Environments

- Logically, provide a "load package" command
- Spack vs. modules:
 - Spack can create TCL or Imod modules
 - Spack can provide its own "environment views" outside of modules
- All these boil down to setting environment variables

Hacking the package stack

- C++:
 - Maintain a "env.sh" file loading appropriate modules
 - Do development there, but be aware that env changes machine to machine
- Python:
 - Create a poetry project to use for its virtual environment.
 - `cd <project>; poetry shell`
 - Keep working scripts / gist-s there.
- Spack:
 - Create a spack environment (`spack env create; spack env activate; spack install ...`)
 - Note also: `spack build-env <project name> bash` (sets CXXFLAGS, etc.)
 - These will load up the environment variables for accessing your installed software.

Intermediate Example: C++ with spack

- <https://github.com/qcscine/sparrow> - semi-empirical quantum chemistry
- git clone <https://github.com/spack/spack>; source spack/share/spack/setup-env.sh; spack compiler find
- spack create <https://github.com/qcscine/sparrow/archive/refs/tags/3.0.0.tar.gz>
 - creates spack/var/spack/repos/builtin/packages/sparrow/package.py
- spack list cereal; spack info boost ~> depends_on("boost@1.65.0:")

Helpful commands:

```
spack dev-build <package> # skip download & build from the current source directory
spack install -u cmake    # download the package & run cmake
spack cd <package>        # change to the directory where spack is working
spack build-env <package> bash # run a shell with env setup to build (and develop)
spack clean               # clears spack's download/build cache
```

https://spack-tutorial.readthedocs.io/en/latest/tutorial_packaging.html

Spack package.py

- `spec = self.spec`
- `spec['mpi'].prefix, spec['mpi'].libs, spec['mpi'].headers`
- <https://spack.readthedocs.io/en/latest/spack.util.html#module-spack.util.prefix>

https://spack.readthedocs.io/en/latest/packaging_guide.html#accessing-dependencies